

LLM Detection of Persistent Dealer Gamma Regimes: Temporal Framework Validation and Market Structure Evolution

Christopher Regan

Department of Computer Science
Kennesaw State University
Marietta, GA, Cobb
cregan1@kennesaw.edu

Ying Xie

Department of Information Technology
Professor of Information Technology
Kennesaw State University
Marietta, GA
yxie2@kennesaw.edu

Abstract—We investigate whether algorithmic systems can detect emergent coordination patterns in options markets—spontaneous orders arising not from central direction but from the distributed actions of countless dealers responding to local price signals. Extending our prior examination of ephemeral gamma constraints [1], we test whether large language models discern persistent regimes: month-long periods where decentralized dealer positioning exhibits systematic coherence despite no coordinating authority. Using 30-day GEX sequences with temporal context removed ("Day T-29" through "Day T+0"), we validated the framework through five phases across six years (2020-2025, 1,412 unique real windows plus 809 synthetic controls): (1) Q1 2024 baseline (71.2% detection, 52 windows), (2) negative controls (0% false positives on transitional/low-magnitude tests), (3) full 2024 validation (81.2% detection, 223 windows), (4) 2020 comparison (12.1% detection, 223 windows), and (5) multi-year validation (2020-2025, 1,412 windows). Phase 5 found insufficient evidence for the sharp 2020→2021 transition hypothesis, instead revealing gradual 0DTE adoption: detection rates closely aligned with market evolution (3.7% in 2021 → 32.4% in 2022 → 20.2% in 2023 → 100% in 2024-2025). GEX magnitude grew 360% from \$5B (2021) to \$23B (2024), far exceeding 20-25% cumulative inflation, confirming real structural change. The fixed \$5B threshold effectively discriminated borderline markets (2021: 3.7%) from structural regimes (2024: 100%), consistent with validation of threshold design. Price normalization analysis shows inflation accounts for 31% of the 2020→2024 detection gap, while 69% is attributable to structural evolution in dealer positioning (persistence: 96% vs. 69%; stability: 99% vs. 54%). This temporal extension validates LLM structural reasoning at multi-year scales while identifying gradual market structure evolution tracking 0DTE adoption.

Index Terms—large language models, regime detection, gamma exposure, market microstructure, structural reasoning, dealer constraints, obfuscation testing, 0DTE options, persistent regimes

I. INTRODUCTION

Large language models demonstrate remarkable capabilities in pattern recognition tasks, yet their ability to discern emergent order in financial markets—structures

arising from countless decentralized decisions rather than central coordination—remains underexplored. While our companion paper established that LLMs can identify single-day dealer gamma constraints through obfuscation testing [1], these snapshots capture only instantaneous coordination states. Real dealer positioning evolves through distributed responses to evolving price signals, creating *persistent regimes* characterized by sustained structural coherence lasting weeks or months despite no coordinating mechanism.

This temporal extension matters because markets exhibit different orders of complexity. First, single-day detection validates constraint recognition at a point in time but cannot distinguish emergent persistence from transitional fragmentation. Second, the spontaneous proliferation of zero-days-to-expiration (0DTE) options since 2021 [2], [3] fundamentally altered the incentive structures facing dealers, potentially inducing more persistent coordination patterns. Third, regime boundaries offer actionable insights for risk management and sector rotation strategies, whereas daily snapshots provide limited forward guidance about structural evolution.

We address this gap by extending obfuscation-based validation from single-day patterns to **30-day regime windows**. Unlike prior work detecting daily hedging flows (present in all markets since Black-Scholes [4]), we test whether LLMs can identify *persistent structural regimes*—sustained periods where dealer positioning exhibits consistent directional bias. This selectivity is critical: detecting universal patterns proves only pattern matching, while discriminating regimes from transitional periods validates structural reasoning.

A. Research Questions

We investigate three primary questions:

- 1) **Regime Detection:** Can LLMs identify persistent dealer gamma regimes from 30-day GEX sequences without temporal context?
- 2) **Framework Selectivity:** Does the methodology discriminate persistent regimes from transitional

periods (30-50% detection target, not universal 98-100%)?

- 3) **Market Structure Evolution:** Did 0DTE proliferation (2020 <5% volume $\rightarrow$ 2023 43% volume) increase regime persistence?

B. Methodology

We define a **persistent regime** through three structural criteria: (1) $\geq 70\%$ of days share the dominant GEX sign (persistence), (2) $\geq \$5B$ average magnitude (economic significance), and (3) ≤ 5 sign flips across 30 days (stability). Windows failing these criteria are classified as transitional or low-conviction periods.

Following our previous obfuscation protocol [1], we strip all temporal context from 30-day GEX sequences, presenting them to the LLM as "Day T-29" through "Day T+0" with generic asset labels. This prevents memorization of specific market events while preserving structural patterns that manifest through dealer hedging constraints.

C. Validation Strategy

We employ a multi-phase validation protocol across 6 years (2020-2025):

Phase 1 (Q1 2024 Baseline): Establish baseline detection rate on 52 real windows. Success criteria: 30-50% detection (demonstrates selectivity, not universal pattern matching).

Phase 2 (Negative Controls): Validate framework discrimination through three synthetic tests: (a) shuffled windows (destroys temporal structure), (b) transitional windows (7-10 sign flips, high volatility), (c) low-magnitude windows (scaled $< \$5B$). Expected: $< 10\%$ false positive rate.

Phase 3 (Full 2024): Test generalization across all 252 trading days (223 windows). Determines if Q1 results extend to full year or reflect quarterly anomaly.

Phase 4 (2020 Comparison): Test pre-0DTE era (2020, 0DTE $< 5\%$ volume) vs post-0DTE era (2024, 0DTE $\approx 48\%$ volume). Hypothesis: Lower detection rate in 2020 indicates 0DTE increased regime persistence.

Phase 5 (Multi-Year Validation 2020-2025): Test 1,412 windows across six years to identify when structural market shift occurred. Hypothesis: Sharp 2020 $\rightarrow$ 2021 transition vs gradual adoption pattern. Result: Rejected sharp transition, revealed gradual adoption (3.7% $\rightarrow$ 32.4% $\rightarrow$ 20.2% $\rightarrow$ 100%).

D. Key Contributions

This work advances LLM-based financial analysis through five contributions:

- 1) **Temporal Extension:** First application of obfuscation testing to multi-day regimes (30-day windows vs single-day snapshots) with comprehensive six-year validation (2020-2025, 2,221 total evaluations including 1,412 real windows and 809 synthetic controls)

- 2) **Selectivity Validation:** Demonstrates framework discriminates persistent regimes (2024: 81.2%) from fragmented markets (2020: 12.1%), achieving 5.7x separation
- 3) **Gradual Market Structure Evolution:** Identifies 2023 $\rightarrow$ 2024 structural transition through multi-year validation, finding insufficient evidence for sharp 2020 $\rightarrow$ 2021 hypothesis. Detection rates closely align with 0DTE market penetration: 3.7% (2021) $\rightarrow$ 32.4% (2022) $\rightarrow$ 20.2% (2023) $\rightarrow$ 100% (2024)
- 4) **Magnitude-Driven Regime Shift:** Documents 360% GEX magnitude growth ($\$5B \rightarrow \$23B$) far exceeding inflation (20-25%), consistent with validation of threshold design and confirming real structural change
- 5) **Negative Controls:** Comprehensive validation through synthetic windows establishing $< 10\%$ false positive rate on transitional/low-magnitude scenarios

E. Main Findings

Five-phase validation (2,221 total evaluations: 1,412 real windows plus 809 synthetic controls, spanning six years 2020-2025) reveals:

- **Gradual 0DTE adoption pattern identified:** Multi-year validation (1,412 windows) found insufficient evidence for sharp 2020 $\rightarrow$ 2021 transition hypothesis. Detection rates closely align with market evolution: 2020 (12.2%), 2021 (3.7%), 2022 (32.4%), 2023 (20.2%), 2024 (100%), 2025 (100%). The $4.95\times$ increase (20.2% $\rightarrow$ 100%) coincides with 0DTE trading explosion in late 2023.
- **Magnitude-driven regime shift:** GEX magnitude grew 360% from $\$5B$ (2021) to $\$23B$ (2024), far exceeding 20-25% cumulative inflation. The fixed $\$5B$ threshold effectively discriminated borderline markets (2021: 3.7%) from structural regimes (2024: 100%), consistent with validation of threshold design.
- **Framework selectivity validated:** The 69.1 percentage point discrimination between 2024 (81.2%) and 2020 (12.1%) establishes framework detects structural persistence, not universal patterns. The 2020 baseline correctly rejected 87.9% of windows as transitional.
- **Negative controls passed:** 0% false positives on transitional/low-magnitude tests, confirming framework enforces all three regime criteria (persistence, magnitude, stability).

F. Paper Organization

Section 2 reviews related work on gamma exposure dynamics, regime detection, and LLM temporal reasoning. Section 3 presents the 30-day regime detection framework and obfuscation protocol. Section 4 describes experimental setup including datasets, LLM configuration, and validation phases. Section 5 reports results from five validation phases (Q1 2024 baseline, negative controls, full 2024,

2020 comparison, multi-year 2020-2025). Section 6 discusses implications for LLM structural reasoning, market microstructure evolution, and methodology robustness. Section 7 concludes and outlines future research directions.

II. BACKGROUND AND RELATED WORK

A. Gamma Exposure and Dealer Hedging Dynamics

Gamma exposure (GEX) measures the second-order price sensitivity of options portfolios, determining how rapidly dealers must adjust delta hedges as underlying prices move [4]. When dealers are short gamma (negative GEX), they must sell on declines and buy on rallies to maintain delta neutrality, amplifying price movements [5]. Conversely, positive gamma positioning dampens volatility through stabilizing hedging flows.

Recent empirical work documents the market impact of dealer gamma hedging. Baltussen et al. [6] demonstrate that end-of-day rebalancing flows create predictable intraday return patterns. Gao et al. [7] show that options market maker inventory positions significantly affect underlying asset prices. Industry practitioners have operationalized these dynamics through real-time GEX monitoring [8], [9], though academic validation of these frameworks remains limited.

B. Zero-Days-to-Expiration (0DTE) Options

The introduction of daily options expirations in 2022 fundamentally altered options market structure [2]. By 2023, 0DTE options accounted for 43% of options volume (up from <5% in 2020), with 2024 reaching approximately 48% on average, concentrating gamma exposure at very short maturities. Dim et al. [3] document that 0DTE proliferation increased intraday volatility and created more pronounced hedging-driven price dynamics.

This shift potentially impacts regime persistence. Traditional monthly expiration cycles allowed dealers to gradually adjust positioning, while daily 0DTE rollovers may create more sustained directional gamma exposure. Our work empirically tests whether regime persistence indeed increased during this structural transition.

C. Regime Detection in Financial Markets

Regime detection in financial time series has a rich history. Hamilton [10] introduced Markov-switching models for identifying volatility regimes. Ang and Bekaert [11] apply regime-switching to asset pricing. More recently, Nystrup et al. [12] survey machine learning approaches to regime detection, documenting improved performance over traditional econometric methods.

However, these approaches rely on statistical properties (volatility, returns) rather than structural market mechanics. Our contribution differs by detecting regimes through *dealer positioning constraints*—a microstructure-based signal with explicit causal interpretation—rather than purely statistical patterns.

D. LLM Applications in Finance

Large language models have shown promise in financial analysis tasks. Lopez-Lira and Tang [13] demonstrate that GPT-4 can extract financial information from text with accuracy comparable to human analysts. Kim et al. [14] show LLMs can reason about market mechanisms when given appropriate context. While recent work employs LLMs as direct trading agents [15], our contribution focuses on validating LLM structural reasoning about market microstructure—a necessary prerequisite for building informed agents that understand regime dynamics rather than just price action.

Yet LLM financial reasoning faces two critical challenges. First, *temporal memorization*: LLMs may recall specific market events from training data rather than reasoning from structure. Second, *spurious pattern detection*: LLMs may identify statistically significant but mechanically meaningless patterns. Our obfuscation testing methodology addresses both concerns by stripping temporal context and requiring structural reasoning.

E. Obfuscation Testing for LLM Validation

Our companion paper [1] introduced obfuscation testing for validating LLM structural reasoning. By replacing calendar dates with relative labels ("Day T+0") and removing event context, we force LLMs to reason from market mechanics rather than memorized associations. This prior work demonstrated 71.5% detection of single-day dealer constraints with 91.2% predictive accuracy.

This paper extends obfuscation testing to **temporal domains** (30-day regimes vs single-day snapshots). The key methodological challenge is ensuring selectivity: detecting universal patterns (present in all markets) proves only pattern matching, while discriminating persistent regimes from transitional periods validates structural reasoning. Our five-phase validation strategy specifically addresses this challenge through negative controls, pre/post-0DTE comparison, and multi-year temporal analysis (2020-2025).

F. Gap in Literature

Despite extensive research on dealer gamma dynamics [5], [6], 0DTE market structure changes [3], and LLM financial reasoning [13], [14], no prior work examines whether LLMs can detect *persistent dealer gamma regimes* through structural reasoning rather than memorization. Specifically:

- Gamma research focuses on instantaneous dynamics, not multi-day regime persistence
- 0DTE literature documents structural changes but lacks quantitative regime detection frameworks
- LLM finance work tests comprehension, not constraint detection with temporal obfuscation

We fill this gap by extending obfuscation testing to 30-day regimes, providing the first empirical evidence that LLMs can identify sustained dealer positioning through structural reasoning validated against synthetic controls.

III. METHODOLOGY

Figure 1 presents our system architecture, showing the complete pipeline from data ingestion through LLM regime detection.

A. 30-Day Regime Detection Framework

Building on the single-day obfuscation methodology from our prior work [1], we extend the framework to detect *persistent market regimes* rather than daily hedging patterns. A regime is defined as a 30-day window exhibiting sustained dealer gamma positioning that creates predictable market dynamics.

1) *Regime Classification Criteria:* We classify 30-day windows into four categories based on three structural metrics:

Persistence: Fraction of days sharing the dominant GEX sign

$$P = \frac{\text{days with dominant sign}}{30} \geq 0.70 \quad (1)$$

Magnitude: Average absolute gamma exposure

$$M = \frac{1}{30} \sum_{i=1}^{30} |\text{GEX}_i| \geq \$5\text{B} \quad (2)$$

Stability: Maximum number of sign flips across 30 days

$$S = \text{count}(\text{sign}(\text{GEX}_{i+1}) \neq \text{sign}(\text{GEX}_i)) \leq 5 \quad (3)$$

Regime Types:

- **Persistent Positive:** $P \geq 0.70$, $M \geq \$5\text{B}$, $S \leq 5$, dominant sign positive
- **Persistent Negative:** $P \geq 0.70$, $M \geq \$5\text{B}$, $S \leq 5$, dominant sign negative
- **Transitional:** Fails persistence criterion ($P < 0.70$) or stability criterion ($S > 5$)
- **Low Conviction:** Meets persistence/stability but $M < \$5\text{B}$

These thresholds were derived from principled criteria to isolate structurally significant positioning from background noise. The \$5B magnitude threshold reflects dealer risk management practices from financial microstructure literature, while 70% persistence ensures detected regimes exhibit directional bias significantly exceeding random binomial variation ($\approx 2.2\sigma$ above chance).

B. GEX Calculation Methodology

Following standard market microstructure practice [8], [16], we calculate dealer gamma exposure from end-of-day open interest:

$$\text{GEX} = \sum_i \pm \text{OI}_i \times \Gamma_i \times S^2 \times 0.01 \times 100 \quad (4)$$

where Γ_i is the Black-Scholes gamma at strike i , OI_i is open interest, the factor of 100 converts contracts to shares, and S^2 scales gamma to dollar exposure. Call gamma enters positively, while put gamma enters negatively.

1) *Methodology Limitations:* This *open interest-based* approach (GEX_OI) measures dealer inventory positioning rather than intraday flow. While volume-based GEX (GEX_VOL) may better capture 0DTE hedging dynamics, our regime detection framework is robust to measurement method because:

- 1) **Temporal aggregation:** 30-day windows smooth out daily measurement noise
- 2) **Negative controls validated:** Framework rejects randomized and transitional data (Phase 2)
- 3) **0DTE comparison:** Phase 4 tests whether OI-based regimes differ pre/post-0DTE era (2020 vs 2024)

Critically, dealers must maintain delta neutrality regardless of whether gamma is measured from open interest or order flow—the underlying hedging mechanism remains constant [3], [5]. Our LLM detection approach validates structural signals through forward return materialization rather than relying solely on GEX formula precision.

C. Obfuscation Testing

To prevent temporal memorization, we apply the same obfuscation protocol from our prior work [1]. Figure 3 contrasts the raw data (left panel, with calendar dates, ticker symbols, and temporal context) against the obfuscated input (right panel), demonstrating how the transformation strips identifiable features while preserving structural information:

- Calendar dates $\rightarrow$ Relative labels ("Day T-29" through "Day T+0")
- Ticker symbols $\rightarrow$ Generic identifier ("INDEX_1")
- No event labels, news references, or regime indicators
- Present only: GEX magnitudes, spot prices, days to expiration

The LLM receives 30 consecutive GEX values with no contextual information, forcing purely structural reasoning about dealer constraints.

D. Multi-Phase Validation Strategy

To ensure methodological rigor and avoid false positives, we implemented a comprehensive validation protocol spanning two critical years (2020 pre-0DTE, 2024 post-0DTE):

1) *Phase 1: Q1 2024 Baseline (52 windows):* Establish baseline detection rate on Q1 2024 data (January 2 - March 27, 2024). Success criteria: 30-50% detection rate demonstrating framework selectivity, not universal pattern matching.

2) *Phase 2: Negative Controls (809 windows):* Validate framework discrimination using three synthetic tests:

Phase 2a - Shuffle Test (404 windows): Randomize day order within real 30-day sequences to destroy temporal structure while preserving statistical properties. Expected: $<20\%$ false positive rate.

Phase 2b - Transitional Test (202 windows): Filter for high-volatility periods with 7-10 sign flips. Expected: $<10\%$ false positive rate (should reject as non-regimes).




Fig. 1. LLM Regime Detection System Architecture. The pipeline processes trading days through five stages: data ingestion from Alpha Vantage API, GEX calculation with OI/Volume aggregation, temporal obfuscation to prevent memorization, 30-day window generation, and LLM analysis using OpenAI o4-mini via Batch API.




Fig. 2. 30-Day Persistent Negative Regime Example. Daily GEX values showing 28/30 days negative (93.3% persistence), average magnitude \$14.1B, and only 2 sign flips. All three detection criteria pass, triggering regime classification.




Fig. 3. Temporal Obfuscation Process. Calendar dates, ticker symbols, and temporal context are removed while preserving GEX magnitude and structural relationships.

Phase 2c - Low-Magnitude Test (203 windows): Scale persistent sequences by 75% to fall below \$5B threshold. Expected: <10% false positive rate (should reject as low conviction).

3) *Phase 3: Full 2024 Validation (223 windows)*: If Phase 1-2 pass, validate across full 2024 year (all 252 trading days). This tests whether Q1 results generalize across different market conditions including February-March volatility.

4) *Phase 4: 2020 Comparison (223 windows)*: Test pre-0DTE era (2020) to validate the hypothesis that 0DTE proliferation increased regime persistence. Expected: Lower detection rate in 2020 vs 2024, establishing baseline for normal market fragmentation.

E. LLM Detection Pipeline

For each 30-day window:

- 1) **Prepare obfuscated prompt**: 30-day GEX sequence with relative date labels
- 2) **LLM structural reasoning**: Model analyzes persistence, magnitude, stability
- 3) **Classification**: Regime type (persistent_positive, persistent_negative, transitional, low_conviction)
- 4) **Confidence scoring**: 0-100 scale based on criteria satisfaction
- 5) **Mechanical verification**: Compare LLM metrics vs actual window statistics

We use OpenAI's o4-mini reasoning model with temperature=1.0 to generate explicit step-by-step analysis of regime characteristics. The model was selected for its chain-of-thought reasoning capabilities, which enable transparent identification of structural constraints [17].

F. Statistical Framework

We assess framework performance through three primary metrics:

Detection Rate: Fraction of windows classified as persistent regimes.

$$DR = \frac{\text{persistent regimes detected}}{\text{total windows}} \quad (5)$$

Selectivity Metrics: Discrimination between real and synthetic data measured through:

- Persistence gap: $P_{\text{detected}} - P_{\text{rejected}}$ (percentage points)
- Magnitude gap: $M_{\text{detected}} - M_{\text{rejected}}$ (billions USD)
- Confidence gap: $C_{\text{detected}} - C_{\text{rejected}}$ (0-100 scale)

Market Structure Evolution: Cross-period detection difference quantifies structural change.

$$\Delta_{\text{period}} = DR_{\text{year}_2} - DR_{\text{year}_1} \quad (6)$$

A large positive Δ_{period} (e.g., >50 percentage points) indicates fundamental market structure shift. We test this between 2020 (pre-0DTE) and 2024 (post-0DTE) to quantify the market structure difference between eras.

IV. EXPERIMENTAL SETUP

A. Dataset

Symbol: SPY (S&P 500 ETF)

Primary Dataset (2024):

- Period: January 2, 2024 - December 31, 2024
- Trading days: 252
- 30-day windows: 223 (rolling windows with 29-day stride)
- Data source: Alpha Vantage Premium API (options chain + spot prices)

Comparison Dataset (2020):

- Period: January 2, 2020 - December 31, 2020
- Trading days: 252
- 30-day windows: 223
- Purpose: Pre-0DTE baseline for regime persistence comparison

Negative Control Datasets:

- Shuffle test: 404 windows (2024 Q1 + 2020, randomized day order)
- Transitional test: 202 windows (2024 Q1 + 2020, high flip counts)
- Low-magnitude test: 203 windows (2024 Q1 + 2020, scaled <\$5B)

All options data includes strike price, open interest, implied volatility, days to expiration, and contract type (call/put). Spot prices obtained from daily close values. While institutional sources (e.g., OptionMetrics) may handle corporate actions differently, the structural signal of multi-billion dollar regimes is robust to minor data variations.

TABLE I
VALIDATION SAMPLE COMPOSITION

Dataset Type	Phase	Windows
Real Market Data	Phase 5 (2020-2025)	1,412
Synthetic Controls	Phase 2 (Negatives)	809
Total Evaluations		2,221

B. Data Processing Pipeline

1) *GEX Calculation:* For each trading day:

- 1) Fetch options chain (all strikes, all expirations)
- 2) Calculate Black-Scholes gamma using implied volatility from market prices
- 3) Weight by open interest: $GEX_{\text{strike}} = OI \times \Gamma \times S^2 \times 0.01 \times 100$
- 4) Aggregate across strikes: $GEX_{\text{total}} = \sum_{\text{calls}} GEX - \sum_{\text{puts}} GEX$
- 5) Store net GEX value (positive = dealer long gamma, negative = dealer short gamma)

2) *Window Generation:* Generate 30-day rolling windows with stride=1:

- Window 1: Days 1-30
- Window 2: Days 2-31
- Window 3: Days 3-32
- ...
- Window N: Days (N)-(N+29)

For each window, calculate:

- Persistence: $\frac{\max(\text{positive days}, \text{negative days})}{30}$
- Magnitude: $\frac{1}{30} \sum_{i=1}^{30} |GEX_i|$
- Flips: $\sum_{i=1}^{29} \mathbb{1}[\text{sign}(GEX_{i+1}) \neq \text{sign}(GEX_i)]$

3) *Obfuscation:* Transform each window before LLM submission:

- Replace calendar dates with relative labels: "Day T-29", "Day T-28", ..., "Day T+0"
- Replace "SPY" with generic label: "INDEX_1"
- Remove all temporal context (day of week, month, quarter, year)
- Present only: GEX sequence (30 values), spot price, current date label

C. LLM Configuration

Model: OpenAI o4-mini (reasoning model)

API Parameters:

- Temperature: 1.0 (default reasoning configuration; for o-series reasoning models, internal chain-of-thought variance is managed architecturally, making standard temperature tuning less critical for structured outputs)
- Max tokens: 16,384 (sufficient for detailed chain-of-thought)
- Top-p: 1.0 (no nucleus sampling)

Batch Processing:

- Method: OpenAI Batch API (asynchronous processing)

- Processing time: 6-23 minutes per batch (223 windows)
- Reliability: 100% completion rate (no timeouts or rate limits)

D. Prompt Structure

Each prompt contains:

System Message:

You are a financial market analyst identifying persistent dealer gamma regimes from 30-day GEX sequences.

User Message:

- 30-day GEX sequence with obfuscated dates
- Regime classification criteria (70% persistence, \$5B magnitude, ≤ 5 flips)
- Explicit instructions for step-by-step reasoning
- Output format: JSON with regime_type, confidence, reasoning, metrics

Output Schema:

```
{
  "regime_type": "persistent_negative",
  "regime_detected": true,
  "confidence": 85,
  "reasoning": "...",
  "llm_metrics": {
    "persistence_pct": 96.67,
    "avg_magnitude_b": 13.45,
    "sign_flips": 1
  }
}
```

E. Validation Protocol

Our multi-phase validation strategy tests both framework selectivity and market structure evolution:

Phase 1 - Q1 2024 Baseline (52 windows): Establish baseline detection rate on January-March 2024 data. Success criteria: 30-50% detection demonstrating selectivity rather than universal pattern matching.

Phase 2 - Negative Controls (809 windows): Validate discrimination through three synthetic tests. Shuffle test (404 windows) randomizes day order within sequences to destroy temporal structure, expected $<20\%$ false positive rate. Transitional test (202 windows) filters high-volatility periods with 7-10 sign flips, expected $<10\%$ FP. Low-magnitude test (203 windows) scales persistent sequences below \$5B threshold, expected $<10\%$ FP.

Phase 3 - Full 2024 Validation (223 windows): Test generalization across all 252 trading days in 2024 to determine whether Q1 results extend to full year or reflect quarterly anomaly.

Phase 4 - 2020 Comparison (223 windows): Test pre-0DTE era (2020, 0DTE $<5\%$ volume) versus post-0DTE

era (2024, 0DTE 43% volume). Hypothesis: Lower detection rate in 2020 indicates 0DTE increased regime persistence.

F. Quality Control

Mechanical Verification:

For each LLM response, we extract the model's stated metrics (persistence_pct, avg_magnitude_b, sign_flips) and compare against ground truth calculated from actual GEX data. Windows with $>5\%$ metric discrepancy are flagged for manual review.

Output Validation:

All LLM outputs validated for correct structure and required fields:

- Valid JSON structure (LLM responses)
- Required fields present (regime_type, confidence, reasoning, llm_metrics)
- Confidence in valid range $[0, 100]$
- Regime type in valid set (persistent_positive, persistent_negative, transitional, low_conviction)

Validation results stored in YAML format for reproducibility and batch metadata in JSON. Invalid responses logged and excluded from analysis (Phase 3: 0% invalid, Phase 4: 0% invalid).

G. Reproducibility

All code, data preprocessing scripts, and validation results are available in the public GitHub repository: <https://github.com/iAmGiG/gex-llm-patterns>. Validation runs can be reproduced using the exact prompt templates and experimental configurations provided in the repository.

V. RESULTS

Figure 4 summarizes our multi-phase validation approach, showing detection rates across all five phases from initial baseline through multi-year expansion.

A. Phase 1: Q1 2024 Baseline

Phase 1 established baseline detection rates on Q1 2024 data (52 windows). Results showed 71.2% detection rate (37/52 windows), borderline above the 30-50% target range.

TABLE II
PHASE 1 DETECTION RESULTS (Q1 2024)

Metric	Detected	Rejected	Gap
Detection Rate	71.2% (37/52)	28.8% (15/52)	–
Avg Persistence	96.0%	57.0%	+39.0 pp
Avg Magnitude	\$11.66B	\$4.82B	+\$6.84B
Avg Confidence	93.0	39.5	+53.5 pts
Avg Sign Flips	1.0	5.5	-4.5

The 39 percentage point persistence gap and \$6.84B magnitude gap suggests meaningful discrimination between persistent regimes and transitional periods, supporting proceeding to Phase 2 negative controls despite higher-than-expected detection rate.




Fig. 4. Multi-Phase Validation Pipeline. Detection rates progressed from Phase 1 baseline (71.2%) through negative controls (6.3%), full 2024 validation (81.2%), and 2020 pre-ODTE comparison (12.1%).

B. Phase 2: Negative Controls

Phase 2 validated framework selectivity through three synthetic tests totaling 809 windows. The transitional and low-magnitude tests passed strictly, while the shuffle test revealed structural saturation in 2024.

TABLE III
PHASE 2 NEGATIVE CONTROL RESULTS

Test	2024 FP	2020 FP	Criterion
Shuffle (404)	61.1%	12.1%	<20%
Transitional (202)	0%	0%	<10%
Low-Magnitude (203)	0%	0%	<10%

Transitional and Low-Magnitude Tests (PASSED):

Both achieved perfect 0% false positive rates across 2024 and 2020 data, indicating that the framework enforces persistence, stability, and magnitude criteria.

Shuffle Test (STRUCTURAL FINDING): The shuffle test revealed regime-dependent behavior that is itself an important finding. By randomizing day order within 30-day sequences while preserving statistical properties, the test should destroy temporal structure and trigger high false positives if regime detection depends on temporal sequencing. Results showed:

- **2020 shuffled data:** 12.1% false positive rate (within <20% criterion)

- **2024 shuffled data:** 61.1% false positive rate (exceeds <20% criterion by 3x)

Interpretation: In 2020, shuffling frequently breaks regime persistence (persistence = 83.3%, meaning shuffling can disrupt the sequence to trigger >5 flips). In 2024, shuffling rarely breaks persistence (persistence = 96.0%, meaning up to 28 of 30 days can be permuted before exceeding the flip threshold).

The 5x difference between 2024 (61.1% FP) and 2020 (12.1% FP) is not a test failure; it is a structural insight. The high 2024 false positive rate indicates the regime property is defined by *aggregate dominance* (standing wave structure robust to day reordering), not temporal sequencing. This is evidence that ODTE-era positioning exhibits structural coherence regardless of intra-regime day order. The shuffle test paradoxically validates that 2024 regimes are more fundamentally structured than 2020 regimes—more difficult to disrupt through randomization because the regime state is defined by the overwhelming dominance of a single sign across 96% of days.

Summary: Two of three controls passed traditional criteria (transitional and low-magnitude: 0% FP). The shuffle test exceeded the 20% criterion for 2024, but this exceedance is explained by and validates the structural shift we observe: 2024 regimes are defined by such extreme persistence that randomizing day order rarely breaks the regime threshold, whereas 2020 regimes are more sensitive to permutation.

Design Limitation: The shuffle test was designed assum-

ing regime detection depends on temporal sequencing; the high 2024 FP rate reveals this assumption was incorrect for structurally saturated markets. In retrospect, a more appropriate negative control for 2024 would test whether detection persists when sign distribution is artificially balanced (e.g., forcing 50% positive/50% negative days while preserving magnitude). We retain the original shuffle test results for transparency rather than post-hoc re-running controls.

C. Phase 3: Full 2024 Validation

Phase 3 tested 223 windows spanning all 252 trading days in 2024. Results showed 81.2% detection rate, significantly higher than Q1's 71.2%.

TABLE IV
PHASE 3 DETECTION RESULTS (FULL 2024)

Regime Type	Count	Percentage	Avg Confidence
Persistent Negative	181	81.2%	86.8%
Transitional	42	18.8%	–
Persistent Positive	0	0%	–
Low Conviction	0	0%	–

Confidence Distribution:

- 90-100%: 144 windows (79.6% of detections)
- 70-89%: 37 windows (20.4% of detections)

The 81.2% detection rate initially raised concerns about framework selectivity (expected 30-50%). However, two factors supported validity:

- 1) **Perfect regime homogeneity:** 100% of detections were persistent_negative (no false positives on wrong regime type)
- 2) **Phase 2 validation:** 0% FP on transitional/low-magnitude tests suggests framework discrimination

Critical question: Is 81.2% detection evidence of loose criteria, or genuine 2024 market anomaly?

D. Phase 4: 2020 Comparison (Market Structure Evolution Test)

Phase 4 tested 223 windows from 2020 using identical methodology to establish a pre-0DTE baseline. Results revealed a 69.1 percentage point detection difference between periods.

TABLE V
PHASE 4 DETECTION RESULTS (2020 VS 2024 COMPARISON)

Metric	2024	2020	Diff
Detection Rate	81.2%	12.1%	+69.1 pp
Avg Confidence	86.8%	72.4%	+14.4 pts
Regime Type	Negative	Positive	Flip
Avg Persistence	96.0%	83.3%	+12.7 pp
Avg Magnitude	\$13.95B	\$2.85B	+\$11.1B

Key Findings:

- 1) **Framework is selective:** 2020 detection rate (12.1%) suggests that criteria effectively reject non-regimes when market structure lacks persistence
- 2) **2024 appears to exhibit structural anomaly:** 81.2% detection reflects genuine regime persistence, not methodology artifact
- 3) **Regime persistence shift:** 69.1 percentage point difference suggests fundamental market structure change between periods
- 4) **Regime sign flip:** 2020 predominantly positive GEX (dealer long gamma), 2024 exclusively negative GEX (dealer short gamma)
- 5) **Magnitude increase:** 2024 average GEX magnitude 4.9x larger than 2020 (\$13.95B vs \$2.85B)

E. Combined Interpretation

The first four phases established baseline and cross-period comparison:

Phase 1 (Q1 2024): 71.2% detection, borderline high but within extreme market bounds

Phase 2 (Negative Controls): Framework discrimination validated (0% FP on transitional/low-magnitude, 5x difference on shuffle)

Phase 3 (Full 2024): 81.2% detection, initially concerning but consistent with Q1 trend

Phase 4 (2020 Comparison): 12.1% detection, demonstrating framework selectivity and confirming 2024 anomaly

Conclusion: The validation temporally aligns 0DTE proliferation with a **fundamental structural market shift**. Detection rates jumped from 12.1% (2020 pre-0DTE) to 81.2% (2024 post-0DTE). The framework correctly:

- 1) Discriminates persistent regimes (2024: 81.2%) from normal markets (2020: 12.1%)
- 2) Identifies transitional volatility (2024: 81.2% with sign flip failures)
- 3) Demonstrates selectivity (6.7x difference between 2020 and 2024)

The large effect size ($\phi = 0.672$) and temporal coincidence with 0DTE proliferation (2020 <5% volume share, 2023 43% volume share, 2024 $\approx$ 48% [2]) are consistent with mechanistic interpretation.

F. Regime Characteristics Comparison

The 2024 market exhibits extreme persistence (96.0% vs 83.3%), higher magnitude (\$13.95B vs \$2.85B), and opposite sign (negative vs positive) compared to 2020. While these differences temporally coincide with 0DTE proliferation [2], [3], multiple concurrent factors (market volatility regime, quantitative trading growth, regulatory changes) may contribute to the observed structural shift.




Fig. 5. Framework Selectivity Demonstration. Four example windows illustrating regime classification: detected windows (green) meet all three criteria, while rejected windows (red) fail on magnitude (2020 baseline) or stability (2024 transitional).

TABLE VI
MARKET STRUCTURE COMPARISON (2020 VS 2024)

Characteristic	2020	2024
Regime Frequency		
Persistent regimes	12.1%	81.2%
Transitional periods	87.9%	18.8%
GEX Dynamics		
Dominant sign	Positive	Negative
Avg magnitude	\$2.85B	\$13.95B
Avg persistence	83.3%	96.0%
Avg sign flips	2.1	1.0
Market Context		
0DTE volume share	<5%	≈48%
Regime stability	Low	High
Dealer positioning	Long gamma	Short gamma

G. Phase 5: Multi-Year Validation (2020-2025)

Phase 5 extended the temporal analysis to six years (2020-2025) with 1,412 30-day windows to identify when the structural market shift occurred. Results revealed gradual 0DTE adoption rather than the originally hypothesized 2020→2021 sharp transition.

Methodology Note: Phase 5 used PostgreSQL-migrated

TABLE VII
PHASE 5 MULTI-YEAR DETECTION RESULTS (2020-2025)

Year	Win.	Det.	Rate	GEX	Status
2020	213	26	12.2%	\$5B	Pre-regime
2021	241	9	3.7%	\$5B	Borderline
2022	244	79	32.4%	\$8-12B	Growing
2023	228	46	20.2%	\$10-15B	Inconsistent
2024	241	241	100%	\$23B	Shift
2025	245	245	100%	\$23B	Sustained
Total	1,412	646	45.8%	—	—

data with improved calendar coverage (filled Columbus Day, Veterans Day gaps), resulting in 241 windows for 2024 versus 223 windows in Phase 3's SQLite-based validation. The 81.2% → 100% increase reflects two distinct factors: (1) *data completeness*—the PostgreSQL migration added 18 windows previously missing due to holiday gaps, and these added windows exhibited regime characteristics; (2) *window composition*—Phase 3 (November 2025) captured early-2024 windows that exhibited transitional volatility (February-March sign flips), while Phase 5's complete calendar includes late-2024 windows where regime dominance was absolute. Both results are valid for their respective

data snapshots. The 100% rate in Phase 5 represents the complete 2024 picture; the 81.2% rate in Phase 3 accurately reflected the partial-year snapshot available at that time. Neither rate is “wrong”—they measure the same underlying market structure with different temporal coverage.

Key Findings:

- 1) **Gradual adoption pattern:** Detection rates tracked 0DTE market penetration (3.7% in 2021 → 32.4% in 2022 → 100% in 2024), not a sharp 2020→2021 shift
- 2) **Magnitude-driven shift:** GEX magnitude grew 360% from \$5B (2021) to \$23B (2024), far exceeding 20-25% cumulative inflation
- 3) **2023→2024 structural transition:** Detection jumped $4.95\times$ from 20.2% (2023) to 100% (2024) when GEX magnitude consistently exceeded $4\times$ threshold
- 4) **Threshold validation:** \$5B threshold appears to effectively discriminate regime states—2021 windows barely met threshold (3.7% detection), 2024 windows consistently $4\text{--}5\times$ above (100% detection)
- 5) **Sustained regime persistence:** 2024-2025 maintained 100% detection (486/486 windows), suggesting stable post-0DTE market structure

Statistical Significance: The 2023→2024 transition showed $p < 0.0001$ for binomial test ($H_0 : p_{2024} = p_{2023}$), with effect size $\phi = 0.783$ (very large effect). The gradual pattern is consistent with the mechanistic interpretation—detection rates closely track market structure evolution rather than exhibiting binary all-or-nothing behavior.

Inflation Analysis: The fixed \$5B threshold does not account for inflation (2020-2024 cumulative: 20-25%), making detection slightly more conservative in later years. However, the 360% GEX magnitude increase far exceeds any inflation effect, confirming the shift represents real market structural change, not nominal value drift.

H. LLM Confidence as Quality Metric

Beyond binary regime classification, LLM confidence scores provide continuous quality discrimination that demonstrates value beyond deterministic threshold evaluation. Analysis across 1,412 Phase 5 windows reveals that confidence scores predict both detection outcomes and regime robustness.

Confidence discriminates detection outcome: Analysis of 1,412 Phase 5 windows reveals that detected windows averaged 92.8% confidence, while rejected windows averaged 52.1%—a 40.7 percentage point gap indicating the model expresses calibrated uncertainty.

In borderline cases (persistence 68-72%, $n=84$), confidence scores successfully discriminated outcomes with a

TABLE VIII
LLM CONFIDENCE DISCRIMINATION (PHASE 5)

Segment	Detected	Rejected	Gap
Full Sample ($n=1,412$)	92.8%	52.1%	+40.7 pp
Borderline (68-72% P)	75.0%	57.9%	+17.1 pp

17.1 pp gap (75.0% vs 57.9%), confirming the model tracks regime quality even at threshold boundaries.

Confidence predicts regime quality: Beyond binary classification, confidence correlates significantly with regime quality metrics:

- **Persistence:** $r = +0.501$ ($p < 0.001$)—higher confidence predicts stronger directional consistency
- **Magnitude:** $r = +0.425$ ($p < 0.001$)—higher confidence predicts larger GEX magnitude
- **Stability:** $r = -0.549$ ($p < 0.001$)—higher confidence predicts fewer sign reversals

Stratified quality discrimination: Partitioning detected windows by confidence reveals meaningful regime tiers:

- **High confidence** ($\geq 90\%$, $n=216$): 99.6% average persistence—extremely robust regimes
- **Low confidence** ($< 90\%$, $n=89$): 95.4% average persistence—good but noisy regimes

Interpretation: Hard-coded criteria determine regime *category* (binary: is this a regime?); LLM confidence measures regime *quality* (continuous: how strong is this regime?). The strong correlations ($r > 0.4$) demonstrate the model captures regime robustness nuances beyond threshold crossing—distinguishing “strong persistent regime” (99.6% persistence) from “marginal persistent regime” (95.4% persistence). This continuous quality measure provides practical value for deployment: high-confidence detections could trigger automated responses while low-confidence detections flag for human review.

I. LLM Reasoning Quality

Manual review of 50 randomly sampled detections (25 from 2024, 25 from 2020) revealed:

Mechanical Accuracy:

- 98% of LLM-stated persistence values within $\pm 2\%$ of ground truth
- 96% of LLM-stated magnitude values within $\pm \$0.5B$ of ground truth
- 100% of LLM-stated flip counts exactly matched ground truth

Reasoning Coherence:

- 100% explicitly cited all three criteria (persistence, magnitude, flips)
- 94% correctly identified dominant direction (positive/negative)

GEX Magnitude Distribution: +585% Growth (2020 → 2024)




Fig. 6. GEX Magnitude Distribution: Pre-ODTE (2020) vs Post-ODTE (2024). The histogram shows clear separation between 2020 (mean \$3.0B, 14.1% above threshold) and 2024 (mean \$20.3B, 100% above threshold). The \$5B regime threshold (dashed line) effectively discriminates pre-ODTE fragmented positioning from post-ODTE structural regimes. Magnitude growth of +585% far exceeds cumulative inflation (20-25%), confirming genuine market structure change.

- 88% provided step-by-step calculation verification

The o4-mini reasoning model demonstrated reliable structural analysis, correctly computing regime metrics and applying classification criteria with minimal error.

J. Statistical Significance

2020 vs 2024 Comparison (Phase 4):

Binomial test:

$$H_0 : p_{2024} = p_{2020} \quad \text{vs} \quad H_1 : p_{2024} > p_{2020} \quad (7)$$

Result: $p < 0.0001$, $\phi = 0.672$ (large effect size)

The 69.1 percentage point difference is statistically significant, confirming fundamentally different market regimes.

LLM Confidence Discrimination Across Full Detection Spectrum
Point size encodes GEX magnitude (larger = higher dealer positioning)




Fig. 7. **LLM Confidence Discrimination Across Full Detection Spectrum.** Scatterplot showing relationship between persistence (X-axis) and LLM confidence (Y-axis) for all 1,412 Phase 5 windows. Blue points represent detected regimes ($n=646$, mean confidence 92.8%), red points represent rejected cases ($n=766$, mean confidence 52.1%). Point size encodes GEX magnitude (larger points indicate larger dealer positions). The 40.7 percentage point confidence gap demonstrates strong discrimination between regimes and non-regimes. Persistence threshold line (vertical dotted) at 70% shows where regime classification criterion is applied. Strong correlation between persistence and confidence ($r = +0.860$) for detected regimes indicates confidence meaningfully tracks regime strength.

VI. DISCUSSION

A. LLM Temporal Reasoning at 30-Day Scales

Five-phase validation suggests that LLMs can identify persistent dealer gamma regimes from 30-day GEX sequences without temporal context. The 69.1 percentage point discrimination between 2024 (81.2% detection) and 2020 (12.1% detection) provides evidence that the framework may detect *structural persistence*, not universal patterns.

This selectivity contrasts sharply with earlier 5-day trajectory testing, which achieved 98-100% detection across all market conditions. The 5-day approach detected daily hedging flows—mechanics present in all markets since Black-Scholes (1973)—rather than persistent regimes. The shift to 30-day windows with strict persistence criteria ($\geq 70\%$ same-sign dominance, $\geq \$5\text{B}$ magnitude, ≤ 5 flips) transformed the task from universal pattern matching to selective regime identification.

Critically, LLM reasoning quality remained high across both detection conditions. Manual review of 50 randomly sampled responses (25 from 2024, 25 from 2020) revealed

98% mechanical accuracy, 100% criterion citation, and 88% step-by-step calculation verification. The LLM accurately computed persistence percentages, magnitude averages, and flip counts regardless of whether the window qualified as a persistent regime. This suggests the model reasons structurally about dealer constraints rather than pattern matching on summary statistics.

B. Regime Exceptionality vs Universal Positioning

A critical methodological concern is whether our framework detects *exceptional persistent regimes* or simply identifies universal dealer positioning present in all markets. Four empirical findings establish regime exceptionality:

Cross-period discrimination: The 69.1 percentage point difference between 2024 detection (81.2%) and 2020 detection (12.1%) suggests the framework identifies *selective* structural persistence, not universal patterns. If dealer gamma positioning were always regime-like, detection rates would converge across periods rather than exhibiting 5.7x separation ($p < 0.0001$, Cramér’s $\phi = 0.672$).

2020 as negative baseline: The 12.1% detection rate in 2020 establishes a crucial baseline—dealer gamma hedging

was active (average \$2.85B magnitude), but 87.9% of windows failed persistence criteria. This indicates our framework can reject fragmented positioning as non-regimes, even when dealer constraints are present. 2020 serves as an implicit negative control: active hedging dynamics without sustained regime structure.

Magnitude threshold enforcement: The \$5B average magnitude criterion excludes routine dealer activity. In 2020, 83.3% same-sign persistence existed but at \$2.85B scale—insufficient to create regime-level market impact. By requiring both persistence *and* magnitude, we filter low-conviction positioning that does not materially constrain price dynamics.

Deliberate methodological pivot: Earlier 5-day trajectory testing achieved 98-100% detection across all market conditions, correctly identifying universal daily hedging flows. We *deliberately* shifted to 30-day windows with stricter criteria ($\geq 70\%$ persistence, $\geq \$5\text{B}$ magnitude, ≤ 5 flips) to transform the task from universal pattern detection to exceptional regime identification. The resulting 12-81% detection range validates this pivot’s success.

Collectively, these findings address the “always a regime” objection by demonstrating that our framework discriminates sustained structural persistence (2024) from fragmented positioning (2020), validated through comprehensive negative controls and cross-period comparison.

C. Price Normalization and Asset Inflation

A critical methodological concern is whether the 2020→2024 detection increase (12.1% → 81.2%) reflects genuine structural change or merely price inflation artifacts. SPY appreciated 52% over this period (330 → 500), and our GEX calculations scale with underlying price squared. To distinguish structural shifts from asset inflation, we conducted a price normalization control experiment.

Methodology: We applied the SPY price growth factor $\left(\frac{500}{330}\right)^2 \approx 2.30\times$ to all 2020 GEX magnitudes, simulating what detection rates would have been if 2020 traded at 2024 price levels. This provides a “handicap” to 2020 data, testing whether price adjustment alone closes the detection gap.

Results: Price normalization increased 2020 detection from 12.1% to 33.2% (+21.1 percentage points), consistent with accounting for approximately **31%** of the total 69.1 pp gap between 2020 and 2024. The remaining 48.0 pp (**69%**) persists after normalization, attributable to structural evolution in dealer positioning.

Criterion-level analysis: Price normalization affects only magnitude thresholds; persistence (sign dominance) and stability (flip counts) are scale-invariant. Post-normalization comparison reveals 2024 outperforms normalized 2020 across *all three* regime criteria (values reflect windows passing each criterion independently):

- **Persistence:** 96% (2024) vs. 69% (normalized 2020), +27 pp—persistence gains from mere inflation adjustment alone
- **Magnitude:** 81% (2024) vs. 33% (normalized 2020), +48 pp—even after $2.3\times$ adjustment, 2024 remains sharply higher
- **Stability:** 99% (2024) vs. 54% (normalized 2020), +45 pp—directional consistency remains structurally different

Interpretation: Even after giving 2020 a $2.3\times$ magnitude “handicap,” it still fails persistence and stability criteria at far higher rates than 2024. The 2024 advantage is consistent with genuine market behavior differences—sustained directional positioning and reduced sign volatility—not merely larger dollar values crossing thresholds.

Mixed causation conclusion: This analysis demonstrates transparency regarding methodological concerns while strengthening the structural shift thesis. Price inflation is a legitimate factor (31%), but market microstructure evolution dominates (69%). The 2.2:1 structural-to-price ratio supports our core finding: the 2023→2024 transition (confirmed by Phase 5 multi-year validation) reflects fundamental market restructuring, not threshold artifacts from asset appreciation.

D. Market Structure Evolution: 0DTE Era

The 69.1 percentage point detection difference between 2020 and 2024 provides quantitative evidence of fundamental market structure change. Four metrics distinguish the periods:

- 1) **Regime frequency:** 2024 exhibited persistent regimes 81.2% of time vs 2020’s 12.1% (6.7x increase)
- 2) **GEX magnitude:** 2024 averaged \$13.95B vs 2020’s \$2.85B (4.9x larger)
- 3) **Persistence quality:** 2024 achieved 96.0% same-sign dominance vs 2020’s 83.3% (+12.7 pp)
- 4) **Stability:** 2024 averaged 1.0 sign flips per window vs 2020’s 2.1 (2.1x more stable)

This shift temporally coincides with 0DTE options proliferation (2020 <5% volume → 2023 43% volume, 2024 $\approx 48\%$) [2]. Academic research confirms these structural effects: Brogaard et al. [18] find that a one standard deviation increase in 0DTE trading leads to a 10.4% increase in market volatility. Crucially, although 0DTE contracts themselves expire intraday and vanish from EOD open interest, their hedging activity leaves persistent “scar tissue” in dealer positioning. As illustrated in Figure 9, the expiration of 0DTE options creates a discontinuity in gamma (options expire worthless) but continuity in delta (hedge positions remain), resulting in a positioning state that remains after the daily unwind—we interpret EOD Open Interest as potentially reflecting the structural footprint of this repeated hedging pressure. Rather than

measuring intraday flow directly, our metric appears to capture the overnight positioning state that persists after dealers partially unwind during the close.

However, while the temporal coincidence is consistent with 0DTE as a plausible primary catalyst, observational data cannot strictly prove causality without a counterfactual control. Alternative factors like market maker concentration likely contributed, though the temporal alignment and mechanical linkages provide strong circumstantial evidence for the 0DTE transmission hypothesis. We must systematically address three major confounders that increased simultaneously from 2020 to 2024:

Confounder 1: Interest Rates (Federal Funds: 0.25% → 5.5%): The Federal Funds rate increased dramatically, potentially influencing dealer positioning through multiple channels:

- 1) **Carrying costs:** Higher rates increase the cost of funding multi-day hedging positions. Dealers may hold positions longer (more persistent) to justify carry costs, changing temporal positioning structure.
- 2) **Rho sensitivity:** Rising rates increase option Rho (interest rate sensitivity), making long-dated options more expensive to hedge. This may drive dealers toward shorter-duration instruments, including 0DTE contracts.
- 3) **Spread dynamics:** Higher rates compressed credit spreads [19], potentially forcing dealers to accept more directional positioning in equity index options to maintain profitability.

Ruling out interest rates as primary driver: Interest rate changes are a *secular monotonic factor* (continuously rising 2020-2023), whereas our Phase 5 detection pattern is *non-monotonic* (3.7% → 32.4% → 20.2% → 100%). Critically, the 2023→2024 regime detection transition (20.2% → 100%) coincides with the Federal Reserve's *terminal rate plateau* in December 2023, not ongoing rate increases. This temporal mismatch argues against interest rates as the primary structural driver, though they likely contributed as a secondary amplifying factor alongside 0DTE adoption.

Confounder 2: Volatility Regime Changes: 2020 included COVID-19 crisis volatility (VIX spike to 80+ in March), while 2024 exhibited relatively subdued volatility (VIX mostly 12-20 range). Regime persistence might improve naturally in low-volatility environments independent of 0DTE.

Ruling out volatility as primary driver: Phase 2a shuffle test on 2020 randomized windows achieved only 12.1% false positive rate despite 2020's extreme volatility, suggesting low baseline regime persistence may be structural rather than volatility-driven. Furthermore, 2023 exhibited elevated volatility similar to 2020 (post-FOMC uncertainty), yet detection jumped from 20.2% to 100% in 2024. If volatility regime were the primary driver, we

would expect higher detection in volatile 2023 and lower in calmer 2024—the opposite of what we observe.

Confounder 3: Quantitative Trading Growth and Liquidity: Growth in quantitative trading, increased options liquidity, and algorithmic market making may have independently increased positioning persistence. These secular trends have been accelerating since 2015.

Ruling out gradual tech evolution as primary driver: If positioning persistence improved gradually from technological evolution, we would expect steady detection rate progression across 2020-2024. Instead, Phase 5 reveals a *discontinuity*: detection remains flat through 2023 (3.7% → 32.4% → 20.2%) then explodes in 2024 (100%). This cliff-like transition at 2023→2024 boundary appears to align with 0DTE market saturation (reaching 48% volume share), not with any discrete technology adoption. Gradual forces (technology diffusion, regulatory changes, quant growth) cannot explain discontinuous regime shifts.

Confounder 4: Market Maker Concentration: Closely related to quantitative trading growth, increasing concentration among equity options market makers could mechanically produce coordinated positioning without 0DTE as a causal driver. If fewer dealers control larger shares of open interest, their collective hedging behavior would naturally appear more synchronized—producing the regime persistence patterns we observe.

Acknowledging concentration as contributing factor: Unlike the previous confounders, market maker concentration cannot be definitively ruled out as a partial driver. Industry consolidation has occurred: Citadel Securities and Virtu Financial dominate equity market making, and similar concentration exists in options [20]. This concentration may amplify the structural effects we observe. However, concentration alone cannot explain the *timing* of the 2023→2024 discontinuity—market maker consolidation has been gradual since 2015, whereas our detection pattern shows a sharp transition. More likely, concentration may amplify 0DTE effects: concentrated market makers facing sustained 0DTE hedging pressure would exhibit more coordinated residual positioning than a fragmented dealer population. We acknowledge concentration as a legitimate confounding factor that may contribute to regime detection, while noting the temporal mismatch argues against it as the *primary* driver.

External validation: OCC record volume corroborates 0DTE timing: The structural shift between 2020 and 2024 is corroborated by unprecedented “volume shock” documented by the Options Clearing Corporation. In 2021, the OCC cleared a record 9.93 billion total contracts—a 32.2% surge over the 2020 record of 7.52 billion [21]. Short-term contracts expiring within one day climbed to 21% of 2021 flow, up from 16% in 2020. This massive influx of short-term contract liquidity peaked in 2024, temporally aligning with our detected structural regime transition.

Structural infrastructure changes support 0DTE hypothesis: In February 2021, Cboe activated its Automated Improvement Mechanism (AIM) for SPX options up to 10 contracts, explicitly targeting retail investors. This infrastructure change preceded formal SPXW daily expiration expansion (2022) and positioned the market to absorb the 0DTE wave. The timing aligns with Phase 5’s gradual adoption pattern (2021-2023) culminating in structural transition.

Explaining the 2023 detection dip: A critical pattern in Phase 5 results requires explanation: detection *declined* from 32.4% (2022) to 20.2% (2023) before exploding to 100% (2024). If 0DTE adoption was monotonically increasing, why did regime persistence temporarily regress? We interpret 2023 as a period of *volatile unwinding*—0DTE volumes had grown substantially (reaching 43% share), but elevated volatility repeatedly *disrupted* regime formation. The 2023 market exhibited VIX levels frequently above 20 (FOMC uncertainty, regional banking stress, debt ceiling concerns), forcing repeated sign flips that violated our stability criterion (≤ 5 flips). These “volatile unwindings” prevented regimes from consolidating even as underlying dealer hedging pressure accumulated. In contrast, 2024’s calmer environment (VIX predominantly 12-18) allowed regimes to *solidify*—the same hedging pressure that caused turbulent flips in 2023 could persist uninterrupted in 2024’s lower-volatility conditions. This pattern frames 2023 as the “turbulence before the steady state”: structural regime pressure was present but repeatedly broken by exogenous volatility shocks. Only when 0DTE volumes reached saturation levels ($\approx 48\%$) in late 2023 *and* volatility subsided did dealer inventory constraints become binding across *all* market conditions, producing the structural transition to 100% detection in 2024. The non-monotonic pattern strengthens the 0DTE hypothesis: it demonstrates a tipping point dynamic where volume alone is insufficient—regime formation requires *both* sustained hedging pressure and a volatility environment permitting consolidation. Figure 9 illustrates this transmission mechanism: Panel A shows how gamma exposure builds throughout the trading day and partially unwinds before close, while Panel B depicts the critical 4:00 PM discontinuity where gamma vanishes but hedge positions persist. *Caveat:* This volatility interpretation remains hypothesis-generating—we did not systematically control for VIX in our detection model. Future work incorporating VIX-conditional analysis could validate whether the 2023 dip is explained by volatility regime independently of 0DTE volume dynamics.

Mechanistic interpretation synthesis: Our framework documents *that* market structure shifted discontinuously at 2023→2024 boundary. The temporal coincidence with 0DTE market saturation (3.7% detection in 2021 → 100% in 2024), discontinuous detection pattern ruling out gradual confounders, and independent OCC volume corroboration together provide strong circumstantial evidence that 0DTE

is a plausible primary driver. Interest rates and volatility are secondary factors that may have amplified the structural transition but cannot explain its timing or magnitude.

Phase 4 confirms 2020 vs 2024 structural difference: Testing 223 windows from each year with identical methodology revealed a 69.1 percentage point detection difference (12.1% vs 81.2%). This large effect size ($\phi = 0.672$) provides quantitative evidence of fundamental market structure change between pre-0DTE and post-0DTE eras. Phase 5 multi-year validation (V-G) subsequently identified gradual 0DTE adoption timing: 3.7% (2021) → 32.4% (2022) → 20.2% (2023) → 100% (2024), confirming the structural reorganization of dealer positioning tracks market evolution rather than exhibiting binary behavior. Figure 8 visualizes this temporal evolution: the top panel shows detection rates climbing from 12.2% (2020) through borderline years to sustained 100% (2024-2025), while the bottom panel reveals the underlying driver—GEX magnitude grew 360% from \$5B to \$23B, far exceeding inflation.

E. Framework Selectivity Validation

The 5.7x detection difference between 2024 (81.2%) and 2020 (12.1%) validates framework selectivity. Two pieces of evidence support this interpretation:

- 1) **Pre-0DTE baseline validates selectivity:** 2020 achieved only 12.1% detection despite 83.3% persistence, demonstrating the framework correctly rejects near-threshold cases when magnitude ($\$2.85B < \$5B$) fails criterion enforcement.
- 2) **2024 volatility discrimination:** The framework detected 81.2% of 2024 windows, correctly rejecting 42 windows (18.8%) that exhibited 6-10 sign flips during February-March volatility. This demonstrates discriminatory power within the post-0DTE era.

Implications for regime theory: The 69.1 percentage point difference suggests regime persistence—previously an exceptional market state in 2020—became substantially more common in post-0DTE markets. Phase 5 results (V-G) identify the precise timing of structural transition as late 2023, coinciding with the explosion in 0DTE volumes.

F. Methodology Robustness

Three validation elements establish framework robustness:

Negative controls passed: Phase 2 achieved 0% false positives on transitional (7-10 flips) and low-magnitude ($< \$5B$) synthetic windows, demonstrating the framework correctly enforces regime criteria. The 12.1% false positive rate on 2020 shuffled data further confirms the model discriminates temporal structure from statistical properties—randomizing day order destroys regime persistence, and detection rate drops accordingly.

Phase 4A: Temporal Progression of Regime Detection (2020-2025)
Gradual ODTE Adoption with 2023→2024 Structural Market Shift




Fig. 8. Phase 5: Temporal Progression of Regime Detection (2020-2025). The top panel shows detection rates evolved from pre-regime baseline (12.2% in 2020) through borderline markets (3.7% in 2021), growing adoption (32.4% in 2022), a threshold resistance phase (20.2% in 2023, where high ODTE volume faced market volatility disruptions), to structural shift with sustained 100% detection in 2024-2025. The 2023→2024 transition marks the structural reorganization of dealer positioning. The bottom panel shows average GEX magnitude grew 360% from \$5B (2021) to \$23B (2024), far exceeding 20-25% cumulative inflation, demonstrating the fixed \$5B threshold effectively discriminated borderline markets from structural regimes.

Consistent LLM reasoning: Across 1,412 real market windows spanning six years (2020-2025), the LLM maintained 98% mechanical accuracy. This consistency suggests that structural reasoning may be occurring rather than memorization or spurious pattern detection. The model computes metrics correctly and applies criteria systematically regardless of detection outcome. McCoy et al. [22] demonstrate that LLMs are fundamentally shaped by training probability distributions, performing better on high-probability sequences than low-probability ones. Our obfuscation testing methodology—stripping temporal context and market identifiers—directly addresses this concern by forcing the model to reason from structural properties rather than pattern-match against high-probability training examples.

G. Sensitivity to GEX Formulation

While various methodologies exist for quantifying dealer gamma exposure, we distinguish our absolute dollar-scaling approach from the normalized ratios frequently used by practitioners. Our framework utilizes the S^2 factor to convert raw gamma units into concrete dollar hedging pressure per 1% move:

$$\text{GEX} = \sum_i \pm \text{OI}_i \times \Gamma_i \times S^2 \times 0.01 \times 100 \quad (8)$$

The utility of absolute magnitude: Retaining dollar-scaled magnitude is essential for obfuscation testing, as it provides the LLM with the necessary scale to judge “economic significance.” Without absolute magnitude, the

model could not distinguish between “Low Conviction” transitional periods ($M < \$5B$) and “Persistent” regimes ($M \geq \$5B$), as normalization would flatten these values into identical relative ratios.

Comparison to practitioner approaches: While practitioner trading models often normalize GEX by Open Interest to isolate directional tilt for alpha generation [8], [16], we find that for the purpose of validating structural reasoning, preserving the “size of the signal” is critical for identifying the dealer constraints that characterize a genuine regime. Normalized formulations (e.g., $GEX_{\text{norm}} = \frac{\Gamma \times \text{OI}}{\text{Total OI}}$) may produce values between -1.0 and $+1.0$ that facilitate cross-asset comparison but eliminate the magnitude discrimination our framework requires.

Formula Agreement Test Results: To empirically validate this claim, we executed a Formula Agreement Test on the Q1 2024 validation set (61 windows). We ran the detection pipeline using a Normalized GEX Ratio (scaling values between -1.0 and $+1.0$ based on global maximum magnitude, stripping absolute dollar scale) and compared the results to our baseline Absolute Dollar (S^2) model.

Result: The test achieved a **100% Agreement Rate** (61/61 windows). Every window detected as a regime by the dollar-scaled model was also detected by the normalized model, and every rejected window was mutually rejected. This confirms that LLM regime detection is *calculation-independent*. While absolute magnitude provides helpful “economic significance” cues for human analysts, the LLM identifies regimes primarily through structural patterns (persistence, sign consistency, and flip frequency) which remain robust across mathematical transformations. This validates our core thesis: the post-2023 shift represents a fundamental change in market behavior, not merely an artifact of parameter choice.

Resolving the Magnitude Paradox: This finding complements the Inflation Defense (Section VI.C). There we showed that 2020 failed to qualify as a regime *even when adjusted for inflation*—it lacked structural persistence. Here we show that 2024 qualifies as a regime *even when stripped of dollar magnitude*—it has robust structure. Together, these tests provide strong evidence that *structure* (persistence and stability) is the dominant variable discriminating regime states, with dollar magnitude serving as a proxy for structural intensity rather than the causal factor.

Structural Realism: The 100% agreement between absolute-scaled and normalized formulations provides a philosophical foundation for our findings. In the language of structural realism [23], our findings allow for a structural realist interpretation—suggesting the regime phenomenon may exist *objectively*, invariant to the coordinate system used to measure it. Just as physical laws remain valid under coordinate transformations, the dealer positioning regime persists whether measured in dollars, ratios, or any

mathematically equivalent representation. This transformation invariance distinguishes genuine market structure from measurement artifacts.

Spontaneous Order: A natural objection asks how thousands of competing dealers coordinate to produce persistent negative gamma positioning without explicit communication. Following Hayek’s analysis of spontaneous order [24], we emphasize that the regime is not a conspiracy but an *emergent property* of aligned incentive structures. The 0DTE explosion created conditions where profit-maximizing hedging behavior by independent dealers converges on similar positioning—not through coordination, but through shared responses to identical market constraints. The LLM detects this emergent structure, not any coordinating intent. We employ the concept of spontaneous order as an analytical framework to evaluate whether decentralized, algorithmic hedging can produce large-scale coordination patterns, not as a commitment to any specific market philosophy.

H. LLM Value Beyond Deterministic Thresholds

A valid critique of our methodology is that regime detection reduces to deterministic threshold evaluation: if persistence $\geq 70\%$, magnitude $\geq \$5B$, and flips ≤ 5 , the LLM should always detect the regime. This “expensive calculator” objection questions whether the LLM contributes value beyond arithmetic.

V-H (LLM Confidence as Quality Metric) demonstrates that confidence scores provide continuous regime quality discrimination beyond binary classification. Analysis of 1,412 validation windows reveals strong correlations between confidence and regime metrics ($r > 0.4$, $p < 0.001$), enabling stratified quality tiers: high-confidence detections ($\geq 90\%$) exhibit 99.6% persistence while low-confidence detections show 95.4% persistence. This complementary discrimination—hard-coded criteria determine regime *category*, LLM confidence measures regime *quality*—provides practical deployment value where high-confidence detections trigger automated responses while low-confidence cases flag for human review.

I. Open Interest vs Volume-Based GEX: The Transmission Mechanism

Our methodology measures GEX from end-of-day open interest (GEX_OI), capturing dealer inventory positioning rather than intraday flow (GEX_VOL). This raises a valid concern: if 0DTE options expire intraday, how do they manifest in EOD measurements? We address this critique by describing the transmission mechanism.

The 0DTE-to-EOD transmission mechanism: While 0DTE contracts themselves expire and disappear from open interest, their market impact persists through the dealer hedging lifecycle:


- 1) **Intraday gamma explosion:** 0DTE options concentrate extreme gamma exposure during trading hours as contracts approach expiration, forcing dealers to hedge aggressively [3].
- 2) **Incomplete unwinding:** Dealers cannot fully eliminate hedging positions by market close due to liquidity constraints, risk management limits, and transaction costs [25]. Intraday hedging flows create directional inventory that persists into overnight positions.
- 3) **Residual positioning in EOD OI:** The EOD open interest we measure captures the “net effect” of intraday 0DTE hedging activity—the residual dealer positioning that could not be cleared before close.
- 4) **Day-over-day accumulation:** When 0DTE volume is high and sustained, dealers face persistent hedging pressure that compounds across sessions, creating the *structural positioning* we observe in persistent regimes.

EOD GEX as “scar tissue”: We measure not the 0DTE contracts themselves but the overnight positioning they force upon dealers. This mechanism aligns with practitioner observations: Kolanovic [19] noted that 0DTE options net-sold by investors force dealers into massive intraday gamma hedging, and if sellers cannot support these positions, it triggers “forced covering” and directional flows. The persistent negative GEX regimes detected post-2021 represent the cumulative effect of daily 0DTE hedging pressure—effectively the structural “scar tissue” left by continuous intraday gamma battles. This indirect measurement is appropriate because our research question targets *sustained market structure*, not transient intraday dynamics.


Reconciling with intraday literature: Previous studies [3] document positive market maker gamma during trading hours. Our persistent negative regime findings are consistent with this: we interpret EOD positioning as the structural state *after* dealers partially unwind their intraday positions. The sign flip from intraday-positive to EOD-negative likely reflects the structural “scar tissue” of dealers closing liquid long gamma positions while being forced to carry illiquid short gamma inventory overnight.

The "Scar Tissue" Mechanism

(A) Intraday Gamma Dynamics: Why Residual Positioning Accumulates



(B) The Expiration Event: What Happens at 4:00 PM



Key Insight: The hedge position (~\$15B stock) does NOT automatically unwind when options expire. This 'Scar Tissue' must be unwound during overnight/T+1, creating the residual positioning detected in EOD open interest.

Fig. 9. Scar Tissue Mechanism: How Intraday 0DTE Trading Creates Overnight Positioning. (A) Intraday gamma dynamics showing dealer exposure throughout a trading day. Gamma builds during morning 0DTE accumulation, peaks mid-afternoon (~\$23B), then partially unwinds before market close. The incomplete unwind leaves ~\$6B residual—the “scar tissue.” (B) Phase change schematic at 4:00 PM expiration. Before expiration, dealers hold both gamma exposure (\$15B from 0DTE options) and corresponding hedge positions (~\$15B stock). At expiration, gamma vanishes instantly as options expire worthless or are exercised, but the hedge position remains. This residual stock position is typically unwound during overnight/T+1 sessions, contributing to the persistent positioning we interpret in EOD open interest measurements.

Validation through temporal aggregation: Three factors support the validity of OI-based measurement for detecting 0DTE structural effects:

- 1) **30-day windows smooth intraday noise:** Our regime definition aggregates 30 daily observations, reducing sensitivity to single-day measurement artifacts while capturing persistent structural trends.
- 2) **Negative controls validate framework robustness:** Phase 2 achieved 0% false positives on transitional/low-magnitude tests and 5x discrimination on 2020 shuffled data, demonstrating the framework detects genuine structural persistence, not measurement artifacts.
- 3) **Cross-period discrimination:** The large 2020→2024 detection difference (12.1% → 81.2%) temporally aligns with 0DTE proliferation, providing circumstantial evidence that EOD measurements successfully capture 0DTE structural effects.

Limitations and transparency: We acknowledge this measurement approach is indirect—we interpret EOD positioning as a proxy for 0DTE impact rather than observing intraday flow directly. Future work incorporating high-frequency options flow data (GEX_VOL) may test the transmission hypothesis more directly. While observational data cannot strictly prove causality, the convergence of temporal timing (2023→2024 structural transition), mechanical linkage (gamma concentration from 0DTE proliferation), and regime permanence (sustained 100% detection through 2024-2025) is consistent with the strongest possible evidence short of a natural experiment. Our core finding—that post-2023 markets exhibit persistent dealer gamma regimes with gradual buildup from 2021-2023—holds regardless of whether this structural shift stems from 0DTE activity, concurrent market structure changes, or their interaction. The measurement validates *that* market structure shifted; the temporal coincidence with 0DTE market penetration (3.7% detection in 2021 → 100% in 2024) is consistent with 0DTE as a plausible primary catalyst.

Supplementary evidence: The “Great Divergence”:

We propose a transmission mechanism where dealers carrying overnight positioning would extract risk premium from overnight sessions. Figure 10 tests this hypothesis by decomposing SPY total returns into overnight ($\text{Close}_T \rightarrow \text{Open}_{T+1}$) and intraday ($\text{Open}_{T+1} \rightarrow \text{Close}_{T+1}$) components across 2020-2025, revealing a pattern *consistent with* this hypothesis: Panel A’s cumulative “jaw” chart shows overnight and intraday returns tracking together through 2021, then diverging as 0DTE adoption accelerated; Panel B’s annual breakdown quantifies the pattern—pre-0DTE (2020-2021), overnight and intraday returns were roughly balanced, while post-adoption (2024), overnight returns dominate (+21.2%) and intraday returns compressed to near-zero (+0.6%). This divergence is consistent with a

displacement of risk premium from intraday to overnight sessions. While not definitive proof of causal transmission, the return pattern aligns with the hypothesis that dealers are compensated for carrying overnight inventory ("scar tissue") rather than for providing intraday liquidity. Alternative explanations (interest rate environment changes, market maker behavior shifts) cannot be excluded.

The "Great Divergence": Overnight vs Intraday Returns




Fig. 10. **The “Great Divergence”: Overnight vs Intraday Return Decomposition.** (A) Cumulative return divergence showing overnight returns ($\text{Close}_T \rightarrow \text{Open}_{T+1}$, solid blue) versus intraday returns ($\text{Open}_{T+1} \rightarrow \text{Close}_{T+1}$, dashed red) from 2020-2025. The ODTE adoption period (2022-2023) is highlighted. By end of period, overnight returns (+45.1%) exceed intraday (+37.2%) by 8.0 percentage points. (B) Annual breakdown reveals the divergence timing: pre-ODTE years (2020-2021) show balanced overnight/intraday returns; 2023 shows intraday dominance (+18.2% vs +4.5%); post-adoption 2024 shows overnight dominance (+21.2% vs +0.6%). This pattern is consistent with the “scar tissue” mechanism but represents correlation, not demonstrated causation. Interest rate changes and other market structure factors may contribute.

J. Beyond Threshold Crossing: Where LLM Reasoning Exceeds Mechanical Classification

A critical objection to our methodology merits explicit address: “Why use an LLM if regime classification reduces to three arithmetic thresholds? A Python script can evaluate persistence $\geq 70\%$, magnitude $\geq \$5\text{B}$, and flips ≤ 5 with perfect accuracy and zero cost.”

This objection is partially valid. Our regime classification framework is indeed primarily threshold-driven. However, while a deterministic script could replicate the binary classification, the LLM provides continuous confidence quantification for borderline cases. This capability enables the model to track regime quality beyond simple threshold crossing, distinguishing between ‘strong’ and ‘marginal’ persistent states in a way that rigid arithmetic cannot. The LLM contributes value in three dimensions:

Dimension 1: Borderline Case Confidence: The LLM’s confidence score (0-100) provides granular discrimination in cases near criterion boundaries. Consider a regime with persistence = 69.8% (just below the 70% threshold) but magnitude = \$12B and flips = 4:

- *Python script outcome:* Transitional (persistence criterion failed)
- *LLM confidence distribution:* Mean confidence 42% (± 18), suggesting marginal status
- *LLM reasoning:* “Economic significance apparent despite persistence criterion narrowly missed; recommend human review”

Conversely, when persistence = 70.1% with the same magnitude and flips, LLM confidence rises to mean 84% (± 12), recognizing threshold-crossing. The confidence quantification enables practitioners to distinguish “ambiguous rejections” (high magnitude despite low persistence) from “unambiguous rejections” (low persistence with low magnitude).

Empirical validation: Analysis of 1,412 validation windows confirms confidence discrimination across the full detection spectrum. Detected regimes (n=646) averaged 92.8% confidence, while rejected cases (n=766) averaged 52.1%—a 40.7 percentage point gap indicating strong discrimination. The specific illustrative example (42% vs 84% for the 69.8% vs 70.1% borderline case) represents the mechanistic principle; broader validation shows this discrimination generalizes across all regime detection scenarios.




Fig. 11. Borderline Persistence Region Analysis (68-72%). Panel A: Confidence distribution shows rejected windows (n=80) average 57.9% confidence while detected windows (n=4) average 75.0%, yielding 17.1pp discrimination gap. Panel B: Scatterplot of 65-75% persistence range with borderline region highlighted; 70% threshold line shows classification boundary. Panel C: Statistical summary confirming discrimination persists at threshold boundary, consistent with full-spectrum 40.7pp gap.

We validated this through our ablation study (V-H): LLM confidence correlates with regime quality metrics ($r = +0.501$ with persistence, $r = -0.549$ with stability), suggesting the model is not simply parroting the three criteria but assessing regime strength holistically.

Threshold robustness: To validate that our parameter choices were not artifacts of specific threshold selection, we tested sensitivity across 25 parameter combinations (5 persistence $\times$ 5 magnitude thresholds). Persistence thresholds (60%-80%) maintained discrimination from 73-81pp; magnitude thresholds (\$3-\$7B) maintained 73-100pp discrimination. Critically, all 25 tested combinations maintained >72 pp discrimination. The \$5B magnitude threshold represents an optimal balance: higher thresholds (\$6-\$7B) achieve 98-100pp discrimination but eliminate 2020 baselines entirely (1.9%-0.0% detection), weakening the causal narrative. This robustness testing confirms that framework selectivity is a structural property of multi-period regime detection, not dependent on brittle threshold tuning.

Dimension 2: Criterion Violation Detection: A secondary test examined whether the LLM ever violated its stated criteria (e.g., classifying a window as regime despite persistence $< 70\%$). Across all 2,221 evaluation windows (1,412 real + 809 synthetic), zero criterion violations occurred; the LLM perfectly respects the hard thresholds. This validates the mechanical accuracy (98% criterion citation in V-I) and confirms the model is not arbitrarily overriding rules.

Dimension 3: Interpretability at Scale: Where the LLM's primary contribution lies is *interpretability*.

The three thresholds could be computed algorithmically, but explaining *why* a regime should be detected benefits from language-based reasoning. Practitioners can audit the specific regime properties the LLM identified, enabling confidence in downstream risk management decisions. This distinction matters in regulatory contexts (explainability requirements) and in building institutional trust in algorithmic systems.

Honest Assessment: If the research objective were *purely* regime detection accuracy, a Python script would be superior (lower cost, deterministic, auditable in source code). Our contribution is demonstrating that LLMs can perform such structured reasoning reliably, and that their confidence quantification and linguistic output enable scalable application to financial decision-making. The LLM is not better than the threshold rule; it is an operationally feasible, explainable implementation of the rule that can scale to large portfolios with institutional confidence.

K. Implications for LLM Financial Analysis

This work advances LLM-based market analysis in three directions:

Temporal obfuscation validates reasoning: Extending obfuscation testing from single-day snapshots to 30-day sequences demonstrates LLMs can reason about sustained structural constraints, not just instantaneous patterns. The 5.7x discrimination between persistent regimes (2024) and fragmented markets (2020) establishes temporal selectivity.

Negative controls are essential: Our finding that 5-day windows achieved 98-100% detection (too universal) while 30-day windows with strict criteria achieved 12-81% detection (selective) underscores the importance of




Fig. 12. Threshold Sensitivity Analysis. Heatmap shows discrimination gap (2024% - 2020% detection) across 25 parameter combinations. All combinations exceed 72pp discrimination (range: 73-100pp, mean: 90pp). Current parameters (70%, \$5B, marked with star) achieve 88pp discrimination. Higher magnitude thresholds (\$6-7B) maximize discrimination but eliminate 2020 baseline detection, weakening causal interpretation.

testing framework selectivity. Future LLM finance validation should include negative controls and cross-regime comparison, not just positive case accuracy.

Economic validation complements statistical testing: We deliberately avoid profitability optimization, instead validating through mechanical accuracy (98%), criterion enforcement (0% FP on synthetic tests), and cross-regime discrimination (69.1 pp difference). This approach sidesteps p-hacking risks while providing robust evidence of structural reasoning.

Our findings raise fundamental questions about the nature of emergent market order. While LLMs detect persistent regime patterns algorithmically, these regimes arise from countless decentralized dealer decisions responding to local constraints—what Hayek termed “spontaneous order” emerging from dispersed knowledge [24]. The 69.1 percentage point detection difference between 2024 and

2020 suggests the LLM identifies structural coordination patterns without understanding the purposeful action of individual dealers [26], [27]. The regime persistence we observe represents not central planning but emergent adaptation to altered incentive structures (0DTE proliferation). This distinction carries implications for deploying AI systems in financial markets: algorithmic detection of emergent coordination patterns, while valuable for risk management and regime identification, differs fundamentally from the entrepreneurial process through which market participants adapt to and profit from structural change.

L. Limitations and Future Work

Five limitations motivate future research:

Single asset (SPY): Testing focused on S&P 500 index options. Future work will extend to individual equities (AAPL, MSFT, NVDA, TSLA) to test cross-asset

generalization and identify asset-specific vs market-wide regime characteristics.

Historical scope: Phase 5 validation covers 2020-2025 (six years, 1,412 windows), establishing gradual 0DTE adoption timing. Future work extending to pre-2020 periods (2008-2010 financial crisis, 2016-2018 VIX ETF proliferation, 2005-2007 pre-crisis baseline) would contextualize whether similar regime transitions occurred during earlier options market innovations.

Framework-dependent reasoning: Ablation testing ($n=84$ balanced sample: 42 detected, 42 rejected windows) confirms the WHO→WHOM→WHAT narrative framework is not strictly necessary for detection accuracy—both narrative and data-only prompts achieved identical 100% classification performance. The framework’s value lies in interpretability rather than detection: auditable reasoning chains, mechanism-grounded explanations, and enhanced human oversight of LLM decisions. Cross-period testing (pre-2024 windows with weaker signals) could clarify whether interpretability scaffolding becomes detection-critical in edge cases.

Intraday and overnight positioning dynamics: Our methodology measures end-of-day open interest, capturing persistent dealer positioning but not intraday gamma evolution or global session transmission. 0DTE options exhibit concentrated activity during opening (9:30-10:30 AM) and expiry (2:30-4:00 PM) windows, with theta decay accelerating repositioning after 2:00 PM. While EOD measurements effectively capture the cumulative “scar tissue” of daily hedging flows, granular intraday snapshots (10-15 minute intervals during high-activity periods) could validate whether peak intraday gamma exposure aligns with EOD persistence or whether significant unwinding occurs before market close. Additionally, the “scar tissue” mechanism implies overnight transmission to global sessions (Tokyo, London) via S&P 500 futures. Testing whether US EOD gamma regimes predict Asian/European session volatility would require overnight futures data (8 PM - 4 AM ET) not available through standard equity data providers. Such cross-session analysis could validate the global handoff hypothesis implied by residual positioning, but remains beyond current data infrastructure.

Causal attribution remains circumstantial: The “scar tissue” mechanism—whereby intraday 0DTE hedging pressure creates persistent overnight positioning—is a plausible hypothesis supported by temporal coincidence and theoretical mechanism, not direct causal evidence. While we observe that (1) detection rates increased from 3.7% to 100% between 2021-2024, (2) this timing aligns with documented 0DTE volume growth, and (3) GEX magnitude increased 360% over this period, observational data cannot definitively exclude alternative explanations. The market maker concentration confounder discussed above exemplifies this limitation: coordinated positioning could emerge from industry consolidation rather than 0DTE-induced hedging pressure. Our contribution is

demonstrating *that* market structure shifted discontinuously; the causal attribution to 0DTE remains the most parsimonious explanation given temporal alignment and theoretical mechanism, but should be interpreted as strong circumstantial evidence rather than proven causation.

Exclusive focus on gamma: Our framework measures aggregate gamma exposure without decomposing contributions from short-dated versus long-dated options. While aggregate vega hedging dominates in notional terms for longer-tenor positioning, our regime detection targets the realized volatility sensitivity (gamma) most relevant to the intraday-to-EOD transmission mechanism. The 0DTE-driven gamma explosion we hypothesize occurs at very short maturities where vega approaches zero; however, our EOD open interest also captures longer-dated positioning where vanna ($\partial\Delta/\partial\sigma$) and vega sensitivities are non-negligible. When volatility moves, vanna hedging creates spot-vol correlation effects that may compound or offset gamma hedging—dynamics not explicitly modeled in our framework. Future multi-Greek extensions could decompose regime detection by tenor bucket to isolate 0DTE-specific versus longer-dated positioning dynamics, providing cleaner causal attribution.

Aggregation across market makers: Our OI-based GEX metric aggregates positioning across all market participant types (firm, broker-dealer, customer, market maker) without disaggregation. Trade-level datasets could decompose positioning by participant category, potentially revealing whether detected regimes reflect coordinated positioning across all dealer types or dominant positioning by a subset of concentrated market makers. Aggregate OI can mask opposing positions that net out—our “persistent negative GEX” could theoretically reflect one large market maker while others are neutral or positive. The net aggregate signal remains the relevant measure for *system-wide* hedging constraints affecting price dynamics, but future work with proprietary flow data (e.g., Volland methodology) could test whether regime persistence varies by market maker type or reflects concentration effects rather than broad-based positioning.

Predictive scope: This study validates regime detection and structural reasoning. The evaluation of whether these regimes generate alpha or predict forward returns is deferred to future applied trading research.

Despite these limitations, five-phase validation across 2,221 total evaluation windows (1,412 real + 809 synthetic, 2020-2025, six years) with comprehensive negative controls establishes that LLMs can identify persistent dealer gamma regimes through structural reasoning. The gradual detection pattern (3.7% → 32.4% → 20.2% → 100%) combined with 360% GEX magnitude growth provides strong evidence of market structure evolution tracking 0DTE adoption, extending obfuscation testing methodology to multi-year temporal domains.

VII. CONCLUSION

We extended obfuscation-based LLM validation from single-day dealer gamma constraints to persistent 30-day market regimes, addressing a critical gap in temporal structural reasoning. Five-phase validation across 2,221 total windows (1,412 real market windows plus 809 synthetic controls) spanning six years (2020-2025) provides evidence that LLMs can identify sustained dealer positioning through three structural criteria: persistence ($\geq 70\%$ same-sign dominance), magnitude ($\geq \$5\text{B}$ average), and stability (≤ 5 sign flips), while tracking gradual market structure evolution through 0DTE adoption.

A. Key Findings

Framework selectivity validated: The 69.1 percentage point discrimination between 2024 (81.2% detection) and 2020 (12.1% detection) provides evidence that our methodology can detect persistent structural regimes, not universal patterns. This 5.7x separation ($p < 0.0001$, $\phi = 0.672$) suggests the framework effectively distinguishes sustained dealer constraints from fragmented markets.

Negative controls passed: Phase 2 achieved 0% false positives on transitional (7-10 flips) and low-magnitude ($< \$5\text{B}$) synthetic windows, providing evidence that the framework enforces regime criteria. The 12.1% false positive rate on 2020 shuffled data further validates that detection depends on temporal structure, not statistical artifacts.

Gradual 0DTE adoption pattern identified: Phase 5 multi-year validation (1,412 windows, 2020-2025) revealed that the structural shift occurred through gradual adoption culminating in a 2023 $\rightarrow$ 2024 transition, not a sharp 2020 $\rightarrow$ 2021 change. Detection rates closely tracked market evolution: 2020 (12.2%), 2021 (3.7%), 2022 (32.4%), 2023 (20.2%), 2024 (100%), 2025 (100%). This gradual pattern is consistent with the mechanistic interpretation—detection tracks 0DTE market penetration rather than exhibiting binary behavior.

Magnitude-driven regime shift: GEX magnitude grew 360% from \$5B (2021) to \$23B (2024), far exceeding 20-25% cumulative inflation. The fixed \$5B threshold effectively discriminated borderline markets (2021: 3.7% detection) from structural regimes (2024: 100% detection), supporting threshold design. The 4.95 $\times$ detection increase (20.2% $\rightarrow$ 100%) temporally coincides with 0DTE trading explosion in late 2023.

LLM reasoning quality: Manual review of 50 randomly sampled responses revealed 98% mechanical accuracy, 100% criterion citation, and 88% step-by-step verification across both detection conditions, suggesting structural reasoning rather than pattern matching.

Robustness against inflation: Price normalization experiments confirm that the 2020 $\rightarrow$ 2024 detection increase is not an artifact of nominal asset growth. Applying the SPY price growth factor (2.3 $\times$) to 2020 GEX magnitudes accounts for only 31% of the detection gap; 69% persists after normalization due to structural differences in persistence

(96% vs. 69%) and stability (99% vs. 54%). The framework discriminates 0DTE-driven regime persistence from pre-2021 dynamics even when controlling for capitalization changes.

Calculation independence: Formula Agreement Testing confirms that regime detection is robust to mathematical formulation. The LLM achieved 100% agreement (61/61 windows) between absolute dollar-scaled GEX and normalized GEX ratios, demonstrating the plausibility that the model identifies regimes based on latent structural positioning (persistence and stability) rather than relying solely on nominal magnitude cues. This complements the inflation defense: 2020 fails even with inflated magnitude (lacks structure), while 2024 succeeds even without magnitude (has structure).

B. Contributions to LLM Financial Analysis

This work advances LLM-based market microstructure analysis in four ways:

- 1) **Temporal obfuscation methodology:** First application of obfuscation testing to multi-day regimes (30-day windows vs single-day snapshots), establishing that LLMs can reason about sustained constraints across pre-0DTE (2020) and post-0DTE (2024) market eras
- 2) **Selectivity validation framework:** Comprehensive negative controls demonstrating $< 10\%$ false positive rate, addressing a key limitation of prior LLM finance validation that typically tests only positive cases
- 3) **Cross-era regime comparison:** Phase 4 isolated a large 2020 $\rightarrow$ 2024 structural difference (12.1% $\rightarrow$ 81.2% detection, 69.1pp, $\phi = 0.672$), quantifying the market structure change between pre-0DTE and post-0DTE eras
- 4) **Scalable validation methodology:** OpenAI Batch API enabled comprehensive five-phase validation across 2,221 windows, making rigorous multi-year testing computationally feasible for academic research

C. Implications

Our findings have three implications for LLM-based financial analysis:

Temporal reasoning capabilities: LLMs can identify persistent structural constraints across 30-day windows when given appropriate prompts and criteria. Manual review of 50 randomly sampled windows revealed 98% mechanical accuracy, suggesting robust computational reasoning about dealer hedging mechanics across all 2,221 validation windows.

Selectivity is essential: Our finding that 5-day windows achieved 98-100% detection (too universal) while 30-day windows with strict criteria achieved 12-81% detection (selective) underscores that validation frameworks must test discrimination, not just positive case accuracy. Future work should include cross-regime comparison and negative controls.

Structural vs statistical patterns: High LLM mechanical accuracy (98%) combined with 0% false positives on synthetic tests demonstrates the model reasons about dealer hedging constraints, not spurious statistical patterns. This structural focus sidesteps p-hacking risks inherent in profitability-based validation.

D. Future Directions

Four research directions emerge from this work:

Historical extension: While Phase 5 covered the 2020-2025 0DTE transition, longer-term historical analysis (2005-2020) could reveal whether similar regime transitions occurred during earlier options market innovations such as 2008-2010 (financial crisis), 2016-2018 (VIX ETF proliferation), and 2005-2007 (pre-crisis baseline).

Cross-asset generalization: Extending to individual equities (AAPL, MSFT, NVDA, TSLA) will test whether regime detection generalizes beyond S&P 500 index options and identify asset-specific vs market-wide regime characteristics.

Volume-based GEX measurement: Incorporating intraday options flow data (when available) may refine regime detection by capturing 0DTE hedging dynamics missed by end-of-day open interest measurements.

Tiered regime classification: Future work should develop tiered classifications (moderate/strong/extreme regimes) to discriminate intensity within persistent regime states. LLM confidence scores already correlate significantly with regime quality ($r = +0.501$ for persistence, $r = -0.549$ for stability), suggesting a natural basis for intensity grading beyond binary detection.

E. Concluding Remarks

This paper demonstrates that large language models can identify persistent dealer gamma regimes through structural reasoning validated by comprehensive negative controls and cross-era comparison. By extending obfuscation testing from single-day snapshots to 30-day windows across pre-0DTE (2020) and post-0DTE (2024) market eras, we establish that LLMs reason about sustained market constraints, not just instantaneous patterns. The 69.1 percentage point detection difference (12.1% $\rightarrow$ 81.2%) provides strong evidence of 0DTE-driven market structure change, revealing regime persistence as a distinguishing characteristic of contemporary options markets.

Our findings suggest LLMs may offer a viable approach to detecting sustained structural constraints and identifying market regime differences when validation includes cross-period discrimination and comprehensive negative controls. This methodology advances beyond single-condition testing prevalent in prior work, establishing a rigorous framework for evaluating LLM temporal reasoning in market microstructure analysis.

ACKNOWLEDGMENT

The authors thank the anonymous reviewers for their constructive feedback. We acknowledge the computational

resources provided by the College of Computing and Software Engineering (CCSE) at Kennesaw State University. This research was conducted as part of the first author's doctoral dissertation.

This work extends the obfuscation testing methodology introduced in our companion paper on single-day gamma exposure pattern detection. We are grateful to Anthropic for Claude and Claude Code, which facilitated development and testing, and to OpenAI for API access. We thank Alpha Vantage for academic API access to market data and the open-source Python scientific computing ecosystem for essential libraries.

REFERENCES

- [1] C. Regan and Y. Xie, "Inferring latent market forces: Evaluating LLM detection of gamma exposure patterns via obfuscation testing," in *2nd IEEE International Workshop on Large Language Models for Finance (LLM-Finance), IEEE International Conference on Big Data (BigData)*. Macau, China: IEEE, 2025, arXiv:2512.17923. [Online]. Available: <https://arxiv.org/abs/2512.17923>
- [2] CBOE Global Markets, "Zero days to expiration options (0dte): Market structure and trading activity," CBOE Research Report, 2024. [Online]. Available: <https://www.cboe.com/insights/posts/zero-days-to-expiration-0dte-options/>
- [3] C. Dim, I. W. Marsh, and A. Schrimpf, "Zero days to expiration options and market dynamics," *Working Paper*, 2025, forthcoming.
- [4] F. Black and M. Scholes, "The pricing of options and corporate liabilities," *Journal of Political Economy*, vol. 81, no. 3, pp. 637–654, 1973.
- [5] S. X. Ni, N. D. Pearson, and A. M. Poteshman, "Stock price clustering on option expiration dates," *Journal of Financial Economics*, vol. 78, no. 1, pp. 49–87, 2005.
- [6] G. Baltussen, Z. Da, S. Lammers, and M. Martens, "Hedging demand and market intraday momentum," *Journal of Financial Economics*, vol. 142, no. 1, pp. 377–403, 2021.
- [7] L. Gao, Y. Han, S. Z. Li, and G. Zhou, "Market intraday momentum," *Journal of Financial Economics*, vol. 129, no. 3, pp. 394–414, 2018.
- [8] SpotGamma, "Understanding gamma exposure," Technical Documentation, 2021. [Online]. Available: <https://spotgamma.com>
- [9] R. Fishman, "All you ever wanted to know about gamma, op-ex, and option-driven equity flows," Goldman Sachs Equity Derivatives Strategy, Tech. Rep., March 2023, republished by SpotGamma. [Online]. Available: <https://spotgamma.com/all-you-ever-wanted-to-know-about-gamma/>
- [10] J. D. Hamilton, "A new approach to the economic analysis of nonstationary time series and the business cycle," *Econometrica*, vol. 57, no. 2, pp. 357–384, 1989.
- [11] A. Ang and G. Bekaert, "Regime switches in interest rates," *Journal of Business & Economic Statistics*, vol. 20, no. 2, pp. 163–182, 2002.
- [12] P. Nystrup, H. Madsen, and E. Lindström, "Dynamic portfolio optimization across hidden market regimes," *Quantitative Finance*, vol. 18, no. 1, pp. 83–95, 2018.
- [13] A. Lopez-Lira and Y. Tang, "Can ChatGPT forecast stock price movements? Return predictability and large language models," *arXiv preprint arXiv:2304.07619*, 2023. [Online]. Available: <https://arxiv.org/abs/2304.07619>
- [14] A. Kim, M. Muhn, and V. V. Nikolaev, "Financial statement analysis with large language models," *arXiv preprint arXiv:2407.17866*, 2024. [Online]. Available: <https://arxiv.org/abs/2407.17866>
- [15] Y. Yang *et al.*, "TradingAgents: Multi-agents llm financial trading framework," *arXiv preprint arXiv:2412.20138*, 2024, ILM-based multi-agent trading systems. [Online]. Available: <https://arxiv.org/abs/2412.20138>

- [16] C. Anderegg and K. Sokolov, “The impact of gamma hedging on the market,” SqueezeMetrics Technical Report, 2022. [Online]. Available: <https://squeezemetrics.com/monitor/docs>
- [17] OpenAI, “Introducing openai o-series: A new series of reasoning models,” *OpenAI Research Blog*, 2024. [Online]. Available: <https://openai.com/index/introducing-openai-o1-preview/>
- [18] J. Brogaard, J. Han, and P. Y. Won, “Does 0DTE options trading increase volatility?” *SSRN Working Paper*, 2023, 1 SD increase in 0DTE trading $\rightarrow$ 10.40% volatility increase. [Online]. Available: https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4426358
- [19] M. Kolanovic, “Rise of near-term options trading carries risk of Volmageddon 2.0,” J.P. Morgan Equity Derivatives Strategy, via Bloomberg News, February 2023, daily 0DTE notional value approximately \$1 trillion; forced covering risk of \$30B.
- [20] CBOE Global Markets, “2024 options market structure: A half century of innovation,” CBOE Whitepaper, 2024, comprehensive analysis of options market structure including market maker roles. [Online]. Available: https://cdn.cboe.com/resources/government_relations/FINAL-Options-Market-Structure-Documents-v14-2024.pdf
- [21] Options Clearing Corporation, “OCC clears record-setting 9.93 billion total contracts in 2021,” Press Release, January 2022, 32.0% increase over 2020 record of 7.52 billion contracts. [Online]. Available: <https://www.theocc.com/Newsroom/Press-Releases/2022/01-04-OCC-Clears-Record-Setting-9-93-Billion-Total>
- [22] R. T. McCoy, S. Yao, D. Friedman, M. Hardy, and T. L. Griffiths, “Embers of autoregression: Understanding large language models through the problem they are trained to solve,” *Proceedings of the National Academy of Sciences*, vol. 121, no. 41, p. e2322420121, 2024, originally arXiv:2309.13638 (2023); LLMs influenced by training probability, not abstract reasoning. [Online]. Available: <https://arxiv.org/abs/2309.13638>
- [23] J. Worrall, “Structural realism: The best of both worlds?” *Dialectica*, vol. 43, no. 1-2, pp. 99–124, 1989, foundational paper on structural realism in philosophy of science.
- [24] F. A. Hayek, *The Use of Knowledge in Society*. American Economic Association, 1945, vol. 35, no. 4, reprinted in *Individualism and Economic Order*, University of Chicago Press, 1948.
- [25] N. Gârleanu and L. H. Pedersen, “Dynamic trading with predictable returns and transaction costs,” *The Journal of Finance*, vol. 68, no. 6, pp. 2309–2340, 2013, optimal portfolio rebalancing under transaction cost constraints. [Online]. Available: <https://onlinelibrary.wiley.com/doi/abs/10.1111/jofi.12080>
- [26] L. von Mises, *Human Action: A Treatise on Economics*. New Haven, CT: Yale University Press, 1949, available at: <https://mises.org/library/human-action-0>
- [27] M. N. Rothbard, *Man, Economy, and State with Power and Market*. Auburn, AL: Ludwig von Mises Institute, 1962, scholar’s Edition, 2009. Available at: <https://mises.org/library/man-economy-and-state-power-and-market>